

Education	<b>PhD Computer Science</b> <i>University of Edinburgh, 2018-2021</i> <ul style="list-style-type: none"><li>- Studying efficient representation structures and optimisations under the supervision of Professor Michael O'Boyle and Professor Amos Storkey.</li><li>- Organising committee member for 1st Workshop on Emerging Deep Learning Accelerators at HiPEAC 2019. Reviewer for CVPR, ICCV.</li><li>- Open source paper replications, tutorials on Gaussian Processes, Python programming.</li></ul>
	<b>MSc Computer Science (Distinction)</b> <i>University of Edinburgh, 2017-2018</i> <ul style="list-style-type: none"><li>- Taught modules on Machine Learning and Parallel Architectures with a project on accelerating training speeds for neural machine translation models.</li><li>- Thesis on hardware adaptive deep learning for embedded GPUs. Replicated deep learning papers in Python and C++ using PyTorch and Tensorflow.</li></ul>
	<b>BSc Computer Science (1st class)</b> <i>University of Birmingham, 2014-2017</i> <ul style="list-style-type: none"><li>- Specialized in machine learning and compilers. Final year courses on statistics, linear algebra and data analysis with project on using neural networks for enhancing stochastic process models.</li></ul>
Technical Skills	<b>Languages:</b> C, C++, Python, OCaml, Bash, CUDA, OpenCL. <b>Other:</b> Deep Learning, PyTorch, Optimising Compilers, Linear Algebra.
Experience	<b>Data Science Intern</b> <i>Lattice Training, 2017</i> Summer internship analysing athlete profiles and using machine learning to predict and develop areas of weakness in rock climbers.
	<b>Technology &amp; Data Summer Analyst</b> <i>Morgan Stanley, 2016</i> Ten week internship in the Technology & Data group at Morgan Stanley. Worked as a full stack engineer building tools for investigating internal dataflow with Scala.
	<b>Campus Ambassador</b> <i>Morgan Stanley, 2015</i> Voted Campus Ambassador of the Year after promoting Morgan Stanley at university events and careers fairs to fellow students.
Publications	<b>BlockSwap: Fisher-guided Block Substitution for Network Compression.</b> <u>J. Turner</u> , E. Crowley, A. Storkey, M. O'Boyle. <i>To appear at International Conference on Learning Representations, 2020.</i>
	<b>Deep Kernel Transfer in Gaussian Processes for Few-shot Learning.</b> M. Patacchiola, <u>J. Turner</u> , E. Crowley, A. Storkey. <i>Currently under submission at AISTATS.</i>
	<b>Distilling with Performance Enhanced Students.</b> <u>J. Turner</u> , E. Crowley, A. Storkey, M. O'Boyle. <i>ARM Research Summit, 2019.</i>
	<b>A Closer Look at Structured Pruning for Neural Network Compression.</b> E. Crowley, <u>J. Turner</u> , A. Storkey, M. O'Boyle. <i>NeurIPS Workshop on Compact Deep Neural Networks, 2018.</i>
	<b>Characterising Cross-Layer Optimisations for Deep Convolutional Neural Networks.</b> <u>J. Turner</u> , J. Cano, V. Radu, E. Crowley, A. Storkey, M. O'Boyle. <i>IEEE International Symposium on Workload Characterisation, 2018.</i>